

23 Bayesian Simple Linear Regression Model

Regression is one of the most widely used statistical techniques for modeling relationships between variables. Regression is also widely used in prediction. We will now consider a Bayesian treatment of *simple linear regression*. We'll use the following example throughout.

Percent body fat (PBF, total mass of fat divided by total body mass) is an indicator of physical fitness level, but it is difficult to **measure accurately**. Body mass index (BMI), which is easily computed based on height and weight ($\text{BMI} = \frac{\text{weight in kg}}{(\text{height in m})^2}$), is another commonly used measure¹. What is the relationship between BMI and percent body fat? How can we use a person's BMI to estimate their percent body fat?

We'll fit a Bayesian model for predicting percent body fat from BMI using data on a nationally representative sample of 2157 U.S. males from the [National Health and Nutrition Examination Survey \(NHANES\)](#).

Example 23.1. Consider the assumptions of a Bayesian simple linear regression model for predicting percent body fat from BMI.

1. What are the assumptions of the simple linear regression model? (There are many different regression models. We're referring to the one that is usually the first one you see in introductory statistics. Sometimes the assumptions are summarized with the acronym "LINE".)
2. What are the parameters in the model? That is, what do we need to put a prior distribution on (in order to eventually get a posterior distribution given sample data)? What do the parameters represent in context?

¹But BMI has its own problems; see <http://fivethirtyeight.com/features/bmi-is-a-terrible-measure-of-health/>.

3. Suggest a prior distribution. Explain your reasoning.
 4. Describe in full detail how you would simulate the prior predictive distribution of percent body fat for males with a BMI of 20.
 5. Simulate the prior predictive distribution of percent body fat for a range of BMI values. Do the distributions seem reasonable? If not, experiment with different priors until you find one that produces a reasonable prior predictive distribution.
- Consider data on a numerical response variable y and a numerical explanatory (or predictor) variable x . A **simple linear regression model** of y on x assumes²:
 - *Linearity*: the conditional mean of y is *linear* in x ; for some slope β_1 and intercept β_0 :

$$E(y|x) = \beta_0 + \beta_1 x$$
 - *Independence*: $y_1, \dots, y_n$ are independent³
 - *Normality*: for each x , the conditional distribution of y given x is Normal
 - *Equal variance*: the conditional variance of y does not depend on x

$$\text{Var}(y|x) = \sigma^2$$
 - The regression model is *conditional* on x . That is, we make assumptions about and perform inference for the conditional distribution of the response variable y given values of the explanatory variable x . Therefore, even when we technically have a sample of pairs (x_i, y_i) from a bivariate distribution, we will always condition on the values of x and treat them as non-random.

²In a Bayesian context, all of these assumptions are also *conditional* on the parameters β_0, β_1, σ

³In a Bayesian context, we assume conditional independence of $y_1, \dots, y_n$ given the parameters β_0, β_1, σ

- Linearity is an assumption: we assume there is some “true” population *regression line* $\beta_0 + \beta_1 x$ which specifies the conditional *mean* of the response variable for each value of the explanatory variable.
- The above set of assumptions can be stated more compactly as

$$y_i \sim N(\beta_0 + \beta_1 x_i, \sigma), \text{ and } y_1, \dots, y_n \text{ are independent}$$

Equivalently,

$$y_i = \beta_0 + \beta_1 x_i + \varepsilon_i, \text{ where } \varepsilon_i, \dots, \varepsilon_n \text{ are i.i.d. } N(0, \sigma).$$

- The parameters of the above model are β_0 , β_1 , and σ .
- σ_y^2 denotes the variance of the response variable y with respect to its overall mean, while σ^2 denotes the variance of values of the response variable about the true regression line. Thus, $\sigma^2 = (1 - \rho^2)\sigma_y^2$, where ρ is the population correlation between the explanatory and response variables.

Example 23.2. Section 23.1.3 summarizes the sample data. Before fitting a Bayesian model, we’ll examine the usual frequentist results.

1. Find and interpret the slope of the fitted regression line.
2. Find and interpret a 95% confidence interval for the slope of the regression line. (How do you interpret 95% confidence?)
3. Find the equation of the fitted regression line.
4. Find and interpret the fitted value for men with a BMI of 20 kg/m².

5. Find and interpret a 95% confidence interval based on the fitted value for men with a BMI of 20 kg/m².
6. Find and interpret a point estimate of σ .
7. Find and interpret a 95% *prediction* interval for men with a BMI of 20 kg/m².

Example 23.3. Now we'll fit a Bayesian model and find the posterior distribution.

1. Suppose there are just two observations of (BMI, PBF): (15, 20) and (40, 30). How would you compute the likelihood?
2. Use `brms` to fit the model to the full sample data. Run some diagnostic checks; does the MCMC algorithm seem to be working efficiently?

3. Inspect the joint posterior distribution of β_0 and β_1 . What do you notice? Why do you think this happens?

4. Now suppose we replace BMI (x) with its centered version $(x - \bar{x})$. What does the parameter β_0 represent in this model? Suggest a prior distribution for β_0 (considering it in isolation).

5. Use `brms` to fit the model using centered BMI as the explanatory variable. Inspect the joint posterior distribution of β_0 and β_1 . What do you notice? Why do you think this happens?

- It is often recommended to center predictors in regression models.
- Centering can help interpretation (because the intercept is often meaningless in context), and also help with multicollinearity (in polynomial models, or when there are multiple predictors).
- In Bayesian framework, centering can reduce the correlation between parameters and improve efficiency of MCMC methods
- When there are multiple numerical predictors it is also often recommended to standardize them (by subtracting the mean and dividing by one or two standard deviations). In a Bayesian framework, standardizing can help set priors when the predictors are on different scales.

Example 23.4. Now we'll summarize and report conclusions based on the posterior distribution, using the model with centered BMI.

1. Describe the posterior distribution of β_1 . Find and interpret 50%, 80%, and 98% posterior credible intervals for β_1 .

2. What does the parameter $\mu_{20} = \beta_0 + \beta_1(20 - \bar{x})$ represent in context? Explain how you could approximate the posterior distribution of μ_{20} . Then approximate the posterior distribution, describe it, and find and interpret a 98% posterior credible interval for μ_{20} .
3. Describe the posterior distribution of σ . Find and interpret 50%, 80%, and 98% posterior credible intervals for σ .
4. Compare the results of the Bayesian and frequentist analyses. How are they similar? How are they different?

Example 23.5. Now we'll use the fitted model to predict percent body fat.

1. Explain how you could approximate the posterior predictive distribution of percent body fat for men with a BMI of 20 kg/m².
2. Find and interpret a 95% posterior prediction interval for men with a BMI of 20 kg/m².

3. Compare the 95% posterior prediction interval with the corresponding frequentist 95% prediction interval. How are they similar? How are they different?
 4. Suggest how you could perform a posterior predictive check. Then use software to perform the check. Does the model seem reasonable based on the data?
- There are a number of ways the assumptions of the simple linear regression model could be revised.
 - We could assume the mean of y changes with x is a non-linear way (e.g., quadratic).
 - We could assume distributions other than Normal in the likelihood. For example, we could assume a t -distribution if outliers were an issue.
 - We could assume the conditional variance of y also changes with x .
 - We could include multiple explanatory variables x .
 - Each of the above would be a change in the *likelihood*, perhaps introducing additional parameters.

Example 23.6. Now we'll investigate sensitivity of the results to choice of prior. Rerun the model (using centered BMI) but let `brms` choose the prior. Compare the results of the two Bayesian models; are the results very sensitive to the choice of prior?

23.1 Notes

23.1.1 Prior predictive tuning

- We'll start with a prior distribution under which β_0 , β_1 , and σ are independent.

- A quick Google search reveals that for men percent body fat ranges from about 6% to 30%, and BMI ranges from about 15 to 35 kg/m².
- The overall range of percent body fat is about 6% to 30%, but we would expect less variability among men of the same BMI. If we assume percent body fat varies by about 5 percentage points among men of the same BMI, we can choose a prior for σ that has a mean of 5, say $\text{Exponential}(1/5)$ (to reflect $\sigma > 0$).
- We expect a positive association between percent body fat and BMI, so we want a prior for β_1 that places most of its credibility on positive values. Disregarding measurement units, the two variables take values on fairly similar numerical scales (6 to 30, 15 to 35). For the slope β_1 , we might assume a Normal prior with a mean of 1 and a standard deviation of 0.5 (to keep most of the credibility on positive values, reflecting prior belief in positive association).
- The intercept β_0 , in principle, represents the percent body fat for a man with a BMI of 0, so we might assume a Normal prior with a mean of 0 and a standard deviation of 1 (to put most of the prior credibility near 0). (We'll see another way to think about this parameter soon.) In summary, our prior distribution on the parameters $(\beta_0, \beta_1, \sigma)$ assumes
 - β_0, β_1 , and σ are independent
 - $\sigma \sim \text{Exponential}(1/5)$
 - $\beta_1 \sim \text{Normal}(1, 0.5)$
 - $\beta_0 \sim \text{Normal}(0, 1)$

Prior predictive distribution for BMI of 20


```
n_rep = 10000

# value of x
x = 20

# simulate parameters from prior distribution
beta0 = rnorm(n_rep, 0, 1)
beta1 = rnorm(n_rep, 1, 0.5)
sigma = rexp(n_rep, 1 / 5)

# simulate values of percent body fat based on assumptions of regression model
y = rnorm(n_rep, beta0 + beta1 * x, sigma)

hist(
  y,
  xlab = "Prior predicted PBF(%) for BMI of 20 kg/m^2",
  breaks = 100,
  col = bayes_col["prior_predict"]
)
```

Prior predictive distribution for many values of BMI

```
n_rep = 1000


# a range of values of x
x = runif(n_rep, 15, 35)

# simulate parameters from prior distribution
beta0 = rnorm(n_rep, 0, 1)
beta1 = rnorm(n_rep, 1, 0.5)
sigma = rexp(n_rep, 1 / 5)

# simulate values of percent body fat based on assumptions of regression model
y = rnorm(n_rep, beta0 + beta1 * x, sigma)

# scatter plot of prior predicted values
plot(
  x,
  y,
  xlab = "BMI (kg/m^2)",
  ylab = "Prior predicted percent body fat (%)",
```

```
col = bayes_col["prior_predict"]
)
```




What regression lines might look like according to prior distribution

```
# number of lines to plot
n_lines = 100

# blank scatterplot - just set the plot area and scale
plot(
  x,
  y,
  xlab = "BMI (kg/m^2)",
  ylab = "Percent body fat (%)",
  type = "n",
  xlim = c(15, 35),
  ylim = c(0, 30)
)

# plot each of the simulated lines
# using slope and intercept simulated from prior distribution
for (l in 1:n_lines) {
```

```
abline(beta0[1], beta1[1], col = bayes_col["prior"], lwd = 1)
}
```



23.1.2 Prior predictive distribution

After some tuning, our prior distribution on the parameters $(\beta_0, \beta_1, \sigma)$ assumes

- β_0 , β_1 , and σ are independent
- $\sigma \sim \text{Exponential}(1/1.5)$
- $\beta_1 \sim \text{Normal}(0.5, 0.1)$
- $\beta_0 \sim \text{Normal}(2, 2)$

Prior predictive distributions

```
n_rep = 1000

# a range of values of x
x = runif(n_rep, 15, 35)

# simulate parameters from prior distribution
beta0 = rnorm(n_rep, 2, 2)
beta1 = rnorm(n_rep, 0.5, 0.1)
```


```

sigma = rexp(n_rep, 1 / 1.5)

# simulate values of percent body fat based on assumptions of regression model
y = rnorm(n_rep, beta0 + beta1 * x, sigma)

# scatter plot of prior predicted values
plot(
  x,
  y,
  xlab = "BMI (kg/m^2)",
  ylab = "Prior predicted percent body fat (%)",
  col = bayes_col["prior_predict"]
)

```



```

# number of lines to plot
n_lines = 100

# blank scatterplot - just set the plot area and scale
plot(
  x,
  y,
  xlab = "BMI (kg/m^2)",


```

```

ylab = "Percent body fat (%)",
type = "n",
xlim = c(15, 35),
ylim = c(0, 30)
)

# plot each of the simulated lines
# using slope and intercept simulated from prior distribution
for (l in 1:n_lines) {
  abline(beta0[l], beta1[l], col = bayes_col["prior"], lwd = 1)
}

```



23.1.3 Sample data

```
data = read_csv("nhanes_males.csv")
```

Rows: 2157 Columns: 2

-- Column specification -----

Delimiter: ","

dbl (2): BMI, PBF

BMI	PBF
20.65	13.1
21.46	31.9
26.76	24.9
26.41	23.5
15.46	25.2
23.90	7.2

- i Use ``spec()`` to retrieve the full column specification for this data.
- i Specify the column types or set ``show_col_types = FALSE`` to quiet this message.

```
data |> head() |> kbl() |> kable_styling()
```

```
c(summary(data$BMI), SD = sd(data$BMI))
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	SD
13.140000	20.150000	23.800000	24.449907	28.100000	45.610000	5.688115

```
c(summary(data$PBF), SD = sd(data$PBF))
```


Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	SD
2.100000	17.400000	23.500000	23.432313	29.200000	51.000000	8.704317

```
cor(data$BMI, data$PBF)
```

```
[1] 0.3981832
```

```
data_plot = ggplot(data, aes(x = BMI, y = PBF)) +
  geom_point(shape = 1, size = 2) +
  labs(x = "BMI (kg/m^2)", y = "Percent body fat (%)") +
  theme_bw()

data_plot
```



23.1.4 Frequentist results

```
freq = lm(PBF ~ BMI, data)
summary(freq)
```

Call:

```
lm(formula = PBF ~ BMI, data = data)
```

Residuals:

Min	1Q	Median	3Q	Max
-21.2198	-5.5169	0.2299	5.4232	26.0046

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	8.53437	0.75906	11.24	<2e-16 ***
BMI	0.60933	0.03024	20.15	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 7.986 on 2155 degrees of freedom
Multiple R-squared: 0.1585, Adjusted R-squared: 0.1582
F-statistic: 406.1 on 1 and 2155 DF, p-value: < 2.2e-16

Confidence interval for mean response

```
x_new = data.frame(BMI = 20)

predict(freq, newdata = x_new, interval = "confidence", se.fit = TRUE)
```


```
$fit
      fit      lwr      upr
1 20.72087 20.29268 21.14907
```

```
$se.fit
[1] 0.2183474
```

```
$df
[1] 2155
```

```
$residual.scale
[1] 7.986374
```

```
data_plot +
  geom_smooth(
    method = lm,
    col = bayes_col["likelihood"],
    fill = bayes_col["likelihood"]
  )
```

Prediction intervals


```
x_new = data.frame(BMI = 20)

predict(freq, newdata = x_new, interval = "prediction")
```

	fit	lwr	upr
1	20.72087	5.053218	36.38853

```
predictions <- predict(freq, interval = "predict")

data |>
  cbind(predictions) |>
  ggplot(aes(x = BMI, y = PBF)) +
  geom_point(shape = 1, size = 2) +
  geom_ribbon(
    aes(ymin = lwr, ymax = upr),
    fill = bayes_col["likelihood"],
    alpha = 0.4
  ) +
  labs(x = "BMI (kg/m^2)", y = "Percent body fat (%)") +
  theme_bw()
```



23.1.5 Posterior distribution

```
library(brms)
```

Loading required package: Rcpp

Loading 'brms' package (version 2.23.0). Useful instructions can be found by typing `help('brms')`. A more detailed introduction to the package is available through `vignette('brms_overview')`.

Attaching package: 'brms'

The following objects are masked from 'package:tidybayes':

`dstudent_t`, `pstudent_t`, `qstudent_t`, `rstudent_t`

The following object is masked from 'package:stats':

`ar`

```
library(tidybayes)
library(bayesplot)
```

This is bayesplot version 1.14.0

- Online documentation and vignettes at mc-stan.org/bayesplot
- bayesplot theme set to `bayesplot::theme_default()`
 - * Does `_not_` affect other `ggplot2` plots
 - * See `?bayesplot_theme_set` for details on theme setting

Attaching package: 'bayesplot'

The following object is masked from 'package:brms':

`rhat`

```
fit <- brm(
  data = data,
  family = gaussian(),
  PBF ~ 1 + BMI,
  prior = c(
    prior(normal(2, 2), class = Intercept),
    prior(normal(0.5, 0.1), class = b),
    prior(exponential(0.667), class = sigma)
  ),
  iter = 3500,
  warmup = 1000,
  chains = 4,
  refresh = 0
)
```

Compiling Stan program...

Start sampling

```
summary(fit)
```

```
Family: gaussian
Links: mu = identity
Formula: PBF ~ 1 + BMI
Data: data (Number of observations: 2157)
Draws: 4 chains, each with iter = 3500; warmup = 1000; thin = 1;
       total post-warmup draws = 10000
```

Regression Coefficients:


	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	8.60	0.72	7.20	9.99	1.00	10059	7435
BMI	0.60	0.03	0.55	0.66	1.00	10239	7616

Further Distributional Parameters:


	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
sigma	7.98	0.12	7.75	8.22	1.00	10922	8170

Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS and Tail_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).


```
plot(fit)
```



```
pairs(fit)
```




```
mcmc_dens_overlay(fit, pars = vars(b_Intercept, b_BMI, sigma))
```



```
neff_ratio(fit)[c("b_Intercept", "b_BMI", "sigma")]
```

```
b_Intercept    b_BMI    sigma
0.7435396    0.7615754  0.8170061
```

```
neff_ratio(fit)[c("b_Intercept", "b_BMI", "sigma")] |>
  mcmc_neff() +
  yaxis_text(hjust = 0)
```



23.1.6 Posterior distribution, model using centered BMI

```
data = data |>
  mutate(BMI_c = BMI - mean(BMI))

fit_c <- brm(
  data = data,
  family = gaussian(),
  PBF ~ 1 + BMI_c,
  prior = c(
    prior(normal(18, 2), class = Intercept),
    prior(normal(0.5, 0.1), class = b),
    prior(exponential(0.667), class = sigma)
  ),
  iter = 3500,
  warmup = 1000,
  chains = 4,
  refresh = 0
)
```

Compiling Stan program...

Start sampling

```
summary(fit_c)
```

```
Family: gaussian
Links: mu = identity
Formula: PBF ~ 1 + BMI_c
Data: data (Number of observations: 2157)
Draws: 4 chains, each with iter = 3500; warmup = 1000; thin = 1;
       total post-warmup draws = 10000
```

Regression Coefficients:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	23.39	0.17	23.05	23.73	1.00	8456	6517
BMI_c	0.60	0.03	0.54	0.66	1.00	11548	7352


Further Distributional Parameters:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
sigma	7.98	0.12	7.75	8.22	1.00	9201	7129


Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS and Tail_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).

```
color_scheme_set("viridis")
```

```
plot(fit_c)
```


```
pairs(fit_c)
```



```
neff_ratio(fit_c)[c("b_Intercept", "b_BMI_c", "sigma")]
```


```
b_Intercept    b_BMI_c    sigma  
0.6516833    0.7352337    0.7128977
```

```
neff_ratio(fit_c)[c("b_Intercept", "b_BMI_c", "sigma")] |>  
  mcmc_neff() +  
  yaxis_text(hjust = 0)
```



Posterior distribution of regression slope

```
mcmc_intervals(fit_c, pars = c("b_BMI_c"), prob = 0.5, prob_outer = 0.98)
```







```
color_scheme_set("teal")

mcmc_areas(fit_c, pars = c("b_BMI_c"), prob = 0.5, point_est = "median")
mcmc_areas(fit_c, pars = c("b_BMI_c"), prob = 0.8, point_est = "median")
mcmc_areas(fit_c, pars = c("b_BMI_c"), prob = 0.98, point_est = "median")
```

Posterior distribution of σ

```
mcmc_intervals(fit_c, pars = c("sigma"), prob = 0.5, prob_outer = 0.98)
```








```
mcmc_areas(fit_c, pars = c("sigma"), prob = 0.5, point_est = "median")
mcmc_areas(fit_c, pars = c("sigma"), prob = 0.8, point_est = "median")
mcmc_areas(fit_c, pars = c("sigma"), prob = 0.98, point_est = "median")
```

Posterior distribution of mean response


```
posterior_c = fit_c |>
  spread_draws(b_Intercept, b_BMI_c, sigma) |>
  mutate(mu_20 = b_Intercept + b_BMI_c * (20 - mean(data$BMI)))

posterior_c |> head(10) |> kbl() |> kable_styling()
```

```
posterior_c |>
  ggplot(aes(x = mu_20)) +
  stat_halfeye(.width = c(0.80, 0.98), fill = bayes_col["posterior"]) +
  theme_bw()
```



.chain	.iteration	.draw	b_Intercept	b_BMI_c	sigma	mu_20
1	1	1	23.27824	0.5574416	8.039549	20.79768
1	2	2	23.62545	0.6372204	7.904723	20.78987
1	3	3	23.62162	0.5754227	7.939876	21.06105
1	4	4	23.12276	0.6061011	8.060678	20.42566
1	5	5	23.12524	0.6090495	8.042735	20.41503
1	6	6	23.36472	0.6317251	8.041233	20.55360
1	7	7	23.50486	0.5616251	7.953135	21.00568
1	8	8	23.73264	0.5591934	7.967661	21.24428
1	9	9	23.01671	0.6313645	7.905566	20.20720
1	10	10	23.32895	0.5444100	7.950075	20.90638



```
mean(posterior_c$mu_20)
```

```
[1] 20.71779
```

```
sd(posterior_c$mu_20)
```

```
[1] 0.2128688
```

.chain	.iteration	.draw	b_Intercept	b_BMI_c	sigma	mu_20	y_pred_20
1	1	1	23.27824	0.5574416	8.039549	20.79768	13.92848
1	2	2	23.62545	0.6372204	7.904723	20.78987	29.75587
1	3	3	23.62162	0.5754227	7.939876	21.06105	42.44593
1	4	4	23.12276	0.6061011	8.060678	20.42566	22.45721
1	5	5	23.12524	0.6090495	8.042735	20.41503	6.98945
1	6	6	23.36472	0.6317251	8.041233	20.55360	22.91524
1	7	7	23.50486	0.5616251	7.953135	21.00568	29.02020
1	8	8	23.73264	0.5591934	7.967661	21.24428	17.71785
1	9	9	23.01671	0.6313645	7.905566	20.20720	17.87872
1	10	10	23.32895	0.5444100	7.950075	20.90638	6.83985

```
quantile(posterior_c$mu_20, c(0.01, 0.99))
```


```
      1%      99%
20.22745 21.21915
```

23.1.7 Posterior prediction

```
posterior_c = posterior_c |>
  mutate(y_pred_20 = mu_20 + sigma * rnorm(nrow(posterior_c), 0, 1))

posterior_c |> head(10) |> kbl() |> kable_styling()
```

```
ggplot(posterior_c, aes(x = y_pred_20)) +
  geom_histogram(
    aes(y = after_stat(density)),
    col = bayes_col["posterior_predict"],
    fill = "white",
    bins = 100
  ) +
  geom_density(linewidth = 1, col = bayes_col["posterior_predict"]) +
  labs(x = "Posterior predicted PBF(%) for BMI of 20 kg/m^2") +
  theme_bw()
```

```
mean(posterior_c$y_pred_20)
```

```
[1] 20.81751
```

```
sd(posterior_c$y_pred_20)
```

```
[1] 8.063809
```

```
quantile(posterior_c$y_pred_20, c(0.025, 0.975))
```

```
      2.5%      97.5%
4.791813 36.748738
```


```
predictive_interval(
  fit_c,
  newdata = data.frame(BMI_c = 20 - mean(data$BMI)),
  prob = 0.95
)
```

```
      2.5%      97.5%
[1,] 4.722616 36.10427
```

23.1.8 Posterior predictive checking


4 simulated samples

```
data |>
  add_predicted_draws(fit_c, ndraws = 4) %>%
  ggplot(aes(x = BMI_c, y = PBF)) +
  geom_point(aes(y = .prediction, group = .draw), size = 0.2) +
  facet_wrap(~.draw)
```



Posterior predictive check of response

```
color_scheme_set("purple")
pp_check(fit_c, ndraw = 100)
```



`posterior_predict` simulates a sample of size 2157, using the observed x values, for each of the 10000 simulated triples from the posterior distribution, resulting in a 10000 x 2157 matrix.

```
predictions <- posterior_predict(fit_c)
```

```
dim(predictions)
```


```
[1] 10000 2157
```

We can use these predictions to plot posterior prediction intervals for a range of BMI values.

```
ppc_intervals(
  data$PBF,
  yrep = predictions,
  x = data$BMI,
  prob = 0.5,
  prob_outer = 0.95
)
```

Warning: Using `size` aesthetic for lines was deprecated in ggplot2 3.4.0.
i Please use `linewidth` instead.

- i The deprecated feature was likely used in the bayesplot package.
Please report the issue at <https://github.com/stan-dev/bayesplot/issues/>.



23.1.9 Sensitivity: using brms default prior

```
get_prior(data = data, family = gaussian(), PBF ~ 1 + BMI_c)
```

	prior	class	coef	group	resp	dpar	nlpar	lb	ub	tag
	(flat)	b								
	(flat)	b	BMI_c							
student_t(3, 23.5, 8.7)	Intercept									
student_t(3, 0, 8.7)	sigma							0		
source										
default										
(vectorized)										
default										
default										

```
fit_brms_prior <- brm(
  data = data,
  family = gaussian(),
  PBF ~ 1 + BMI_c,
  sample_prior = TRUE,
  iter = 3500,
  warmup = 1000,
  chains = 4,
  refresh = 0
)
```

Compiling Stan program...

Start sampling

```
summary(fit_brms_prior)
```

```
Family: gaussian
Links: mu = identity
Formula: PBF ~ 1 + BMI_c
Data: data (Number of observations: 2157)
Draws: 4 chains, each with iter = 3500; warmup = 1000; thin = 1;
       total post-warmup draws = 10000
```

Regression Coefficients:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	23.43	0.17	23.09	23.76	1.00	9860	7258
BMI_c	0.61	0.03	0.55	0.67	1.00	9646	7651


Further Distributional Parameters:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
sigma	7.99	0.12	7.75	8.23	1.00	10509	7583


Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS and Tail_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).

```
color_scheme_set("viridis")

plot(fit_brms_prior)
```




```
pairs(fit_brms_prior)
```



```
neff_ratio(fit_brms_prior)[c("b_Intercept", "b_BMI_c", "sigma")]
```


b_Intercept	b_BMI_c	sigma
0.7258387	0.7651354	0.7583361

```
neff_ratio(fit_brms_prior)[c("b_Intercept", "b_BMI_c", "sigma")] |>  
  mcmc_neff() +  
  yaxis_text(hjust = 0)
```



```
color_scheme_set("purple")
```

```
pp_check(fit_brms_prior, ndraw = 100)
```




```
predictions_b <- posterior_predict(fit_brms_prior)

dim(predictions_b)
```

```
[1] 10000 2157
```

```
ppc_intervals(
  data$PBF,
  yrep = predictions_b,
  x = data$BMI,
  prob = 0.5,
  prob_outer = 0.95
)
```

Leave-one-out (LOO) cross-validation to compare models.

```
LOO(fit_c, fit_brms_prior)
```

Output of model 'fit_c':

Computed from 10000 by 2157 log-likelihood matrix.

	Estimate	SE
elpd_loo	-7544.2	31.0
p_loo	2.7	0.1
looic	15088.3	62.0

MCSE of elpd_loo is 0.0.

MCSE and ESS estimates assume MCMC draws (r_eff in [0.8, 1.1]).

All Pareto k estimates are good (k < 0.7).
See help('pareto-k-diagnostic') for details.

Output of model 'fit_brms_prior':

Computed from 10000 by 2157 log-likelihood matrix.

	Estimate	SE
elpd_loo	-7544.2	30.9
p_loo	2.8	0.1
looic	15088.5	61.9

MCSE of elpd_loo is 0.0.

MCSE and ESS estimates assume MCMC draws (r_eff in [0.9, 1.1]).

All Pareto k estimates are good ($k < 0.7$).
 See `help('pareto-k-diagnostic')` for details.

Model comparisons:

	elpd_diff	se_diff
fit_c	0.0	0.0
fit_brms_prior	-0.1	0.4